

— CONCEPT DOCUMENT

Most marketing tools start at the post. This one starts at the business.

An AI marketing system that runs a real consultation first, tells a business owner what is actually costing them money, and only then builds the plan, writes the content and runs the spend. Plus a second track that turns the work into paid, supervised training for students.

Small businesses do not have a content problem. They have a system problem.

Every marketing tool on the market opens with the same question: *what do you want to post?* That question assumes the business behind the post is ready to receive the attention. Most of the time, it is not.

A shop owner who spends on ads while messages sit unanswered for a day is not under-marketed. A brand that does not know its margin cannot tell a winning campaign from a slow loss, and will not find out until the money is gone. A business with no way to record where a customer came from cannot repeat anything that works, so every quarter starts from zero.

None of these are content problems, and no content generator will surface them. It will cheerfully produce thirty posts for a business that loses most of its buyers between "how much?" and payment — and then the owner concludes that marketing does not work.

The gap is not between businesses that post and businesses that don't. It is between businesses that have a system and businesses that have a feed.

Why the existing tools cannot close it

Canva, Predis, AdCreative, Buffer, Later, Hootsuite and the rest are genuinely good at what they do. But they are all downstream of the decision that matters. They are production tools sold to people whose problem is diagnosis, and the reason is structural: diagnosis requires knowing what to ask, which is consulting, not software.

The other half of the gap is language and context. An international tool writes Modern Standard Arabic that nobody in a Beirut comment section has ever written — grammatically perfect and socially wrong. It assumes card checkout when the market runs on cash and local wallets. It treats WhatsApp as a support channel when here it is the till.

TO COMPLETE

Add your own evidence here before submitting: how many small businesses you have spoken to, what they said about their current spend and returns, and any letters of intent. Cite every market-size figure to a named, dated source. A number without a source is the fastest way to lose a judging panel.

A consultation, a diagnosis, then a machine that runs the answer.

RankRise Ascend is one continuous system in five stages. You answer questions about your business — not your feed. It scores how ready you are to market, names what is costing you money, builds a 90-day plan that repairs exactly those things, writes the content that plan needs, and then runs, watches and reports on it.

The order is the product. Each stage is only possible because the one before it happened, and each stage visibly answers the one before it. The plan cites the diagnosis. The calendar cites the plan. Nothing is generic, because nothing starts from a blank page.

01 Consult

A structured intake of 34 questions across seven sections: the business, the offer, the customer, the current presence, the system behind the counter, the next 90 days, and the competition. It asks about margin, reply speed, how orders arrive and what breaks when you get busy — the things an agency asks in the first meeting and no software asks at all.

OUTPUT A complete, structured picture of the business.

02 Diagnose

A Marketing Readiness Score out of 100, six scored pillars — offer, brand, channels, conversion, measurement, ad readiness — and findings ranked by what each one costs, every finding carrying its own first fix. Plus a direct verdict on whether this business should be spending on ads at all yet, and a short list of what must be true before it does.

OUTPUT One number the owner remembers, and a ranked list they can act on.

03 Plan

A 90-day plan in three phases, sized to the hours and budget the owner actually reported. Every action names the diagnosis finding it repairs. Phase one stops what leaks, phase two builds the engine, phase three pushes volume through it. With a budget split, a publishing cadence, a first-seven-days checklist and the specific risks that would derail it.

OUTPUT A plan the owner can start on Monday, not a strategy deck.

04

Create

Three or four content pillars so nobody ever faces a blank page, then four weeks of slots at the cadence the plan committed to — each slot carrying a format the business can actually produce and an angle tied to a plan action. One tap turns any slot into a finished post: hook, caption, hashtags, shot list, alt text. In English, Modern Standard Arabic, spoken Lebanese, or the mixed register most brands here actually write in.

OUTPUT A month of work, ready to shoot and publish.

05

Run

Scheduling and publishing, story reminders, performance reporting in plain language, and a competitor watch built on the ad libraries the platforms publish by law — so the owner is told when a rival launches a campaign, and what it appears to be selling. Media-buying plans and boost strategies are prepared by the system and launched by the owner.

OUTPUT The plan actually happens, and the owner learns from it.


A worked example, end to end.

This is the example that ships inside the product so the flow can be walked without a network. The business is invented — it is not a client, and it carries no client's data — but the reasoning is exactly what the engine produces from real answers.

The business

A one-branch specialty coffee shop in Beirut. Roasts in-house. Sells bags at \$14. Around 5,000 followers, posts a few times a month, no ads beyond occasional boosts, orders arrive by walk-in, WhatsApp and Instagram DM. The owner runs marketing between everything else and has one to three hours a week. Budget: under \$200 a month. Asked about margin, the owner answered *"I honestly don't know."*

The diagnosis



"Your coffee is not the problem — your counter is."

The engine found three blockers before it found anything to post about: messages answered "whenever we see it", an unknown margin, and no record of where any customer came from. Its verdict on ad spend was **not yet** — not because \$200 is too little, but because money spent now could not be measured and the traffic it bought could not be answered. It also found one strength worth everything: the roast date on the bag is a specific, checkable claim no competitor on that street can make, and almost nothing the business publishes says it.

The plan that follows

Three phases — *Stop the leak*, *Make the bag look like the coffee*, *Reach past the street* — with every action naming the finding it repairs. Media spend stays at zero until day 61, because the diagnosis said so. The north star is bag sales per month, not followers: bags do not queue behind one espresso machine, they carry the best margin, and they reach people who will never walk down that street.

The content it writes

Twelve slots across four weeks at three posts a week — the cadence the plan sized to one to three hours.
One tap on the first slot, written in spoken Lebanese:

شوف التاريخ عالكيس.

كل كيس عنا فيه تاريخ التحميص. مش تاريخ إنتاج، مش تاريخ صلاحية — تاريخ التحميص.

هيدا الكيس تحمّص اليوم. رح يوصلك وهو لسا ما مرق عليه أسبوعين.

روح شوف الكيس يلي عندك بالبيت هلق. إذا ما في تاريخ تحميص عليه، في سبب.

Shipped with a shot list a person with a phone can follow, hashtags mixing product, city and brand reach, and alt text for accessibility. The working notes stay in English; only what gets published is in the content language.

Four things a competitor cannot copy by shipping a feature.

"AI writes your posts" is not a moat — it is a commodity with better-funded incumbents. These four are where the actual defensibility sits, and they are ordered by how hard they are to copy.

1 · The consultation is productised agency judgement

The scoring rubric is not arithmetic on dropdown values — a weighted sum of menu indices would produce a number that looks precise and means nothing. It is a written judgement framework: what a low margin implies for paid spend, why a dormant account scores below no account, why an unknown margin is itself a low score. That framework comes from running campaigns for real clients, and it improves every time a diagnosis is checked against what the business actually did next. A competitor can copy the questions in an afternoon. They cannot copy the answers.

2 · Language that matches how the market writes

Every international generator returns Modern Standard Arabic. The rules inside Ascend are specific down to individual words — عايز not بدى, دلوقتي not هلق, + عم verb for the present, رح for the future — and they keep English mixed in the way Lebanese brands genuinely write it, rather than sanitising it out. "Write in Lebanese dialect" as an instruction reliably returns Egyptian with a few words swapped. This is market knowledge encoded as rules, not a dropdown.

3 · Built for how business is actually done here

WhatsApp is a primary sales channel, not a support inbox. Cash and local wallets dominate over card checkout. Instagram outweighs almost everything else. Budgets of \$150 a month are the normal case, not a degenerate one — and a business with \$200 needs a different strategy, not a scaled-down version of a \$5,000 one. Tools built elsewhere get all of this wrong by default, and fixing it is not a setting.

4 · It comes from an agency with live clients

RankRise already runs this process by hand for real businesses. That is the difference between "we have an idea for a tool" and "we have a system that works and we are productising it" — and it means every client engagement is both revenue and a test of the engine.

THE RULE THAT PROTECTS ALL OF IT

The engine is forbidden from inventing a metric, a benchmark, a competitor statistic or a case study. Where it has no number it describes the mechanism instead. This is enforced in the system's own instructions, not left to chance — because a fabricated percentage collapses the first time anyone checks, and it would take the product's credibility with it.

RankRise Academy – the work becomes paid training.

The same platform has a second surface. Marketing work that comes out of stage four and five has to be executed by somebody. Students need real experience and cannot get it. Small businesses need execution and cannot afford an agency. Ascend already knows exactly what work needs doing, because it wrote the plan.

How it works

1. **A student enrolls and picks a track** — content, design, video, paid media, copy — and is levelled by a short assessment rather than by what they claim.
2. **The system assigns a real task from a real plan**, scoped to that level: a caption set, a carousel, a short edit, an ad variant. The task comes from a live business's calendar, not from an exercise sheet.
3. **The student submits work with their reasoning attached.**
4. **The AI reviews it against the brief** — what works, what does not, what would be rejected by a client and why — and sets a timed re-attempt where it matters.
5. **A human supervisor signs off** before anything reaches a client. The AI accelerates review; it does not replace judgement, and a student's grade is never decided by a model alone.
6. **The student earns a RankRise certificate** backed by a portfolio of actual published work, and the strongest ones get hired.

Why this is strong, not a side feature

It is a two-sided marketplace

Businesses get execution at a price they can pay. Students get experience nobody else will give them. Each side makes the other more valuable, which is the hardest thing to build and the hardest to displace.

It is a talent pipeline

The agency sees hundreds of students doing graded real work before hiring any of them. That is a recruiting advantage no job posting can match.

It has genuine social weight

Paid, supervised, portfolio-backed training for students in a market with very few routes into the industry. This is real impact, not impact framing.

It closes the loop

Every graded submission is data on what good work looks like for a given brief — which makes the engine that wrote the brief better.

BUILT IN FROM DAY ONE, NOT BOLTED ON

Client data reaching students is a real privacy obligation. The design requires explicit client consent per engagement, an anonymisation mode that strips identifying details from training tasks, a signed undertaking from every student, and supervisor sign-off before any student work reaches a client account. Expect a judge to ask about this; the answer should be ready.

Three customers, one platform.

SEGMENT	WHAT THEY HAVE TODAY	WHAT ASCEND GIVES THEM
The owner-operator Shop, clinic, café, salon, trades	Runs marketing themselves between everything else. A few hours a week at most. Posts when they remember. Spends on boosts without knowing what came back.	A straight answer on what is broken, a plan sized to the hours they really have, and content ready to publish without hiring anyone.
The small marketing team One or two people in-house	Capable but stretched. Spends most of its time producing rather than deciding. No structured way to justify budget upward.	The diagnosis and plan as an argument to management, and the calendar and content as the production load lifted.
The agency Small shops and freelancers	Wins clients on instinct. Onboarding and audits eat unbilled hours. Quality varies by whoever is assigned.	A white-label audit and plan that makes onboarding a repeatable product, plus Academy students as vetted capacity.
The student Academy side	A degree, no portfolio, and no route to real client work. Internships are scarce and mostly unpaid.	Graded real tasks, a supervisor, a portfolio of published work, a certificate, and a route to being hired.

Where we start

Lebanon first — it is where RankRise operates, where the language and payment realities are understood, and where the gap is widest. The same conditions hold across the region and in diaspora-owned businesses, which is the expansion path: Gulf, Egypt, Türkiye, then Arabic and English-speaking diaspora markets. The product does not need rebuilding to travel; the language rules and the payment assumptions are configuration, not architecture.

Five streams, and only one of them is a subscription.

A single subscription line is fragile — it caps at what a small business will pay monthly, and in this market that ceiling is low. These five compound instead, and three of them grow with the customer rather than with the headcount.

STREAM	WHAT IT IS	WHY IT HOLDS
1 · Subscription	Tiered monthly access to the platform. A free diagnosis, a paid tier for plan and content, a higher tier for publishing, competitor watch and multiple brands.	Predictable base. The free diagnosis is the acquisition engine — it delivers real value before asking for anything.
2 · Execution marketplace	A commission on work executed through the Academy: the business pays for the task, the student is paid, the platform takes a margin.	Grows with usage, not with sales headcount, and it is the piece competitors cannot bolt on.
3 · Academy fees	Students or their institutions pay for enrolment, supervision and certification. Universities are a channel, not just a customer.	Counter-cyclical to the business side, and institutionally sticky once a university adopts it.
4 · Media-spend share	A percentage of managed ad spend, for businesses that let the platform prepare and run campaigns.	Scales with the customer's success. The one line that grows without the customer feeling a price rise.
5 · White label	Agencies license the diagnosis and plan under their own brand, per seat or per client audit.	Turns the nearest competitor into a distribution channel.

TO COMPLETE BEFORE SUBMISSION

Set your actual price points per tier, the commission percentage, and your own unit economics — what it costs you to serve one account per month against what that account pays. Judges will ask, and the answer has to be yours, not a placeholder. Price in the currencies and rails your customers actually use: local wallets and cash matter more here than card subscriptions, and the billing design has to reflect that from the start.

Why the free diagnosis is the whole go-to-market

The stage that costs least to deliver is the one that creates most trust, because it tells people things they did not know and sometimes tells them not to spend money. A business that receives an honest "do not run ads yet" has been given something no competitor's free tier offers. That is what earns the second conversation — and the second conversation is where the plan, the content and the execution are sold.

What is built today, honestly.

This is not a concept looking for an engineer. Four of the five stages run end to end against a live model, in a working application, and can be demonstrated on request.

STAGE	WHAT IT DOES	STATE
01 Consult	34-question intake across seven sections, saved as you go, required-answer gating	Built
02 Diagnose	Readiness score, six scored pillars, ranked findings with fixes, ad-spend verdict	Built — live model
03 Plan	90-day plan in three phases, every action tied to a finding, budget split, week-one checklist	Built — live model
04 Create	Content pillars, four-week calendar, one-tap post writing in four language registers	Built — live model
05 Run	Publishing, competitor ad watch, reporting	In progress
Academy	Student tracks, task assignment, AI review, supervisor sign-off	Designed, not started

How it is built

A static front end with a small serverless back end, so it deploys anywhere in minutes and costs almost nothing to run while it is small. The model API key never reaches the browser. Every generating call streams, so a screen never sits blank while the engine thinks — and the progress a user sees is driven by bytes actually received, never by a timer pretending to be progress.

Every generated document is defined as a strict schema and validated before it reaches the screen, so a half-formed result fails loudly rather than rendering as a broken score. And the whole flow works with no API key at all, on a bundled worked example that is clearly marked as a sample — because venue wi-fi fails, and a demo that dies without a network dies on stage.

What is deliberately not automated yet

Publishing directly to Instagram and TikTok requires platform app review, which takes weeks of waiting rather than weeks of work. Until that is granted, the system prepares the post and the owner presses publish. The same applies to spending money: campaigns are built and costed by the system and launched by the owner. Handing a machine unsupervised access to someone's ad budget is a decision that should be earned, not defaulted to.

What happens next, in order.

HORIZON	WHAT SHIPS
Now	Stage 5: scheduling, story reminders, plain-language reporting, and competitor ad watch built on public ad libraries. Owner-launched publishing while platform review is pending.
Next	Accounts and saved brands, so a diagnosis can be re-run in 90 days and scored against the first one. Brand kit memory so every generated asset stays on-brand. Client approval flow — the owner sees, approves, and then it publishes.
Then	Academy v1: student tracks, levelling, task assignment from live calendars, AI review with supervisor sign-off, certificates. Launched with one university partner rather than opened to everyone.
After	Platform publishing approvals granted and automated. Generated visuals and short video from the brand kit. Agency white-label. Regional language packs beyond Lebanese.

Risks, and what we do about them

RISK	RESPONSE
Platform API approval is slow and outside our control	The product is designed to be fully useful without it. Prepared-and-launched is the default; automated publishing is an upgrade, not a dependency. No part of the pitch rests on approval we do not have.
An incumbent adds a "diagnosis" feature	They can copy the questions; the rubric is judgement built from running real campaigns, and it compounds with every engagement. The language and market knowledge are a second wall, and the Academy is a third they cannot build from outside.
The model produces something wrong or invented	Invented metrics and case studies are forbidden in the engine's own instructions; every output is schema-validated; and nothing in the Academy reaches a client without a human sign-off. Where the system has no number it says so.
Client data exposed through the Academy	Explicit consent per engagement, an anonymisation mode for training tasks, signed undertakings from students, and supervisor review before anything is published.
Competitor intelligence crosses a line	Public sources only — the ad libraries platforms publish by law, and what anyone can see from outside an account. No scraping of private data, and the product never implies otherwise.

RISK**RESPONSE**

Small businesses churn after the free diagnosis

The diagnosis is designed to be re-run at 90 days and scored against the first one. Progress against your own previous number is the retention mechanic, and it only exists because stage two exists.

What the funding buys.

The engine works. What money buys is the distance between a working engine and a business: the execution layer, the Academy, and enough runway to prove retention with real paying customers rather than with a demo.

Engineering

Complete stage five, ship accounts and brand memory, and build Academy v1. The largest single line, and the one that turns four working stages into a product people pay for every month.

Model and infrastructure

The per-run cost of the engine at real volume, plus hosting. Small per account, and the one cost that rises directly with usage.

Academy operations

Supervisors, the first university partnership, and the legal and privacy work that has to be right before a single student touches a client account.

Go to market

Getting the free diagnosis in front of enough businesses to learn what converts, and building the agency white-label channel.

TO COMPLETE BEFORE SUBMISSION

Replace this section with your actual figures: the amount requested, the split across these four lines, the runway it produces, and the specific milestones that money reaches. State what you have already invested — the working product is itself evidence of commitment. If the competition asks for projections, build them from your own unit economics and label every assumption. Do not include a number you cannot defend in the room.

What we will have proven by the end of that runway

- That businesses complete a 34-question consultation — the single biggest product risk, and the one a demo cannot answer.
- That the diagnosis changes what an owner does next, measured by re-running it at 90 days.
- That the execution marketplace clears: real tasks, real students, real client acceptance.
- That at least one university will put its name to the Academy certificate.

We are not asking anyone to believe an idea works. Four of the five stages are built and can be walked through right now. The ask is for the distance between that and a business.

